

Conditionals: Advanced WHERE

Complementary SQL: Exercises

May 11, 2025

Contents

1	Practice Meanings, Values, Relations, Advantages, and Unique Uses	4
1.1	Subquery with <code>IN</code>	4
1.2	Subquery with <code>EXISTS</code>	4
1.3	Subquery with <code>ANY</code>	4
1.4	Subquery with <code>ALL</code>	4
1.5	<code>IS DISTINCT FROM</code>	4
1.6	<code>IS NOT DISTINCT FROM</code>	4
2	Practice Disadvantages of All Its Technical Concepts	5
2.1	<code>NOT IN</code> with Subquery Returning <code>NULL</code>	5
2.2	<code>!= ANY</code> Misinterpretation	5
2.3	Performance of <code>IS DISTINCT FROM</code> vs. Standard Operators (Conceptual)	5
2.4	Readability of <code>EXISTS</code> vs. <code>IN</code> for Simple Cases	5
3	Practice Inefficient/Incorrect Alternatives vs. Advanced WHERE Conditions	6
3.1	Inefficient <code>COUNT(*)</code> vs. <code>EXISTS</code> for Existence Check	6
3.2	Verbose/Incorrect <code>NULL</code> Handling vs. <code>IS DISTINCT FROM</code>	6
3.3	Complex <code>NULL</code> -aware Equality vs. <code>IS NOT DISTINCT FROM</code>	6
3.4	Using <code>LEFT JOIN</code> and checking for <code>NULL</code> vs. <code>NOT EXISTS</code>	6
4	Hardcore Problem Combining Previous Concepts	7

Dataset for PostgreSQL

The following SQL code creates and populates the necessary tables for the exercises.

```
1  -- Drop tables if they exist to ensure a clean setup
2  DROP TABLE IF EXISTS employee_projects CASCADE;
3  DROP TABLE IF EXISTS projects CASCADE;
4  DROP TABLE IF EXISTS employees CASCADE;
5  DROP TABLE IF EXISTS departments CASCADE;
6
7  -- Create Departments Table
8  CREATE TABLE departments (
9      dept_id SERIAL PRIMARY KEY,
10     dept_name VARCHAR(50) NOT NULL,
11     location VARCHAR(50),
12     monthly_budget NUMERIC(10,2) NULL
13 );
14
15 -- Create Employees Table
16 CREATE TABLE employees (
17     emp_id SERIAL PRIMARY KEY,
18     emp_name VARCHAR(100) NOT NULL,
19     dept_id INTEGER REFERENCES departments(dept_id),
20     salary NUMERIC(10,2) NOT NULL,
21     manager_id INTEGER REFERENCES employees(emp_id) NULL,
22     performance_rating INTEGER NULL CHECK (performance_rating IS NULL OR
23     performance_rating BETWEEN 1 AND 5),
24     last_bonus NUMERIC(8,2) NULL,
25     hire_date DATE NOT NULL
26 );
27
28 -- Create Projects Table
29 CREATE TABLE projects (
30     proj_id SERIAL PRIMARY KEY,
31     proj_name VARCHAR(100) NOT NULL,
32     lead_emp_id INTEGER REFERENCES employees(emp_id) NULL,
33     budget NUMERIC(12,2),
34     start_date DATE NOT NULL,
35     end_date DATE NULL
36 );
37
38 -- Create Employee_Projects Junction Table
39 CREATE TABLE employee_projects (
40     emp_id INTEGER REFERENCES employees(emp_id),
41     proj_id INTEGER REFERENCES projects(proj_id),
42     role VARCHAR(50),
43     hours_assigned INTEGER NULL,
44     PRIMARY KEY (emp_id, proj_id)
45 );
46
47 -- Populate Departments
48 INSERT INTO departments (dept_name, location, monthly_budget) VALUES
49 ('Human Resources', 'New York', 50000.00),
50 ('Technology', 'San Francisco', 75000.00),
51 ('Sales', 'Chicago', 60000.00),
52 ('Support', 'Austin', 40000.00),
53 ('Research', 'Boston', NULL), -- Budget is NULL
54 ('Operations', 'New York', 50000.00);
55
56 -- Populate Employees
57 -- Top Managers (no manager_id)
58 INSERT INTO employees (emp_name, dept_id, salary, manager_id, performance_rating,
59     last_bonus, hire_date) VALUES
60 ('Alice Wonderland', 1, 90000.00, NULL, 5, 10000.00, '2010-03-15'),
61 ('Bob The Builder', 2, 95000.00, NULL, 4, 8000.00, '2008-07-01');
62
63 -- Other Employees
64 INSERT INTO employees (emp_name, dept_id, salary, manager_id, performance_rating,
65     last_bonus, hire_date) VALUES
66 ('Charlie Brown', 1, 60000.00, 1, 3, 3000.00, '2012-05-20'), -- HR
67 ('Diana Prince', 2, 75000.00, 2, 5, 7000.00, '2015-11-01'), -- Tech
68 ('Edward Scissorhands', 2, 70000.00, 2, 2, NULL, '2016-02-10'), -- Tech, NULL bonus, low
69     rating
```

```

66 ('Fiona Apple', 3, 65000.00, NULL, 4, 5000.00, '2018-08-01'), -- Sales, no manager_id in
    this context
67 ('George Jetson', 3, 55000.00, 6, 3, 2500.00, '2019-01-15'), -- Sales
68 ('Hannah Montana', 4, 50000.00, 1, NULL, 1500.00, '2020-06-01'), -- Support, NULL rating
69 ('Ivan Drago', 4, 48000.00, 8, 2, 1000.00, '2021-03-10'), -- Support
70 ('Julia Child', 5, 80000.00, NULL, 5, NULL, '2011-09-05'), -- Research, NULL bonus
71 ('Kevin McCallister', 1, 58000.00, 1, 4, 2000.00, '2013-07-22'), -- HR
72 ('Laura Palmer', 2, 82000.00, 2, 3, 4000.00, '2014-01-30'), -- Tech
73 ('Michael Knight', 3, 68000.00, 6, 5, 6000.00, '2017-04-11'), -- Sales
74 ('Nancy Drew', 4, 52000.00, 8, 4, NULL, '2019-10-01'), -- Support, NULL bonus
75 ('Oscar Wilde', 5, 78000.00, 10, NULL, 7500.00, '2022-01-20'); -- Research, NULL rating
76
77 UPDATE employees SET manager_id = 1 WHERE emp_name = 'Charlie Brown';
78 UPDATE employees SET manager_id = 1 WHERE emp_name = 'Kevin McCallister';
79 UPDATE employees SET manager_id = 2 WHERE emp_name = 'Diana Prince';
80 UPDATE employees SET manager_id = 2 WHERE emp_name = 'Edward Scissorhands';
81 UPDATE employees SET manager_id = 2 WHERE emp_name = 'Laura Palmer';
82 UPDATE employees SET manager_id = 6 WHERE emp_name = 'George Jetson';
83 UPDATE employees SET manager_id = 6 WHERE emp_name = 'Michael Knight';
84 INSERT INTO employees (emp_name, dept_id, salary, manager_id, performance_rating,
    last_bonus, hire_date) VALUES
85 ('Peter Pan', NULL, 30000.00, NULL, 3, NULL, '2023-01-01'); -- No department, NULL bonus
86
87 -- Populate Projects
88 INSERT INTO projects (proj_name, lead_emp_id, budget, start_date, end_date) VALUES
89 ('Alpha Launch', 4, 150000.00, '2023-01-01', '2023-12-31'), -- Lead: Diana Prince (Tech)
90 ('Beta Test', 5, 80000.00, '2023-03-01', '2023-09-30'), -- Lead: Edward Scissorhands (
    Tech)
91 ('Gamma Initiative', 1, 200000.00, '2022-06-01', NULL), -- Lead: Alice Wonderland (HR)
92 ('Delta Rollout', 13, 120000.00, '2024-02-01', NULL), -- Lead: Michael Knight (Sales)
93 ('Epsilon Research', 10, 90000.00, '2023-05-01', '2024-05-01'), -- Lead: Julia Child (
    Research)
94 ('NoLead Project', NULL, 50000.00, '2023-07-01', NULL); -- NULL lead_emp_id
95
96 -- Populate Employee_Projects
97 INSERT INTO employee_projects (emp_id, proj_id, role, hours_assigned) VALUES
98 (4, 1, 'Developer', 160), -- Diana on Alpha
99 (5, 1, 'QA Engineer', 120), -- Edward on Alpha
100 (12, 1, 'UI Designer', 100), -- Laura on Alpha
101 (5, 2, 'Lead Tester', 150), -- Edward on Beta
102 (9, 2, 'Tester', 80), -- Ivan on Beta
103 (1, 3, 'Project Manager', 200), -- Alice on Gamma
104 (3, 3, 'Coordinator', NULL), -- Charlie on Gamma, NULL hours
105 (11, 3, 'Analyst', 100), -- Kevin on Gamma
106 (13, 4, 'Sales Lead', 180), -- Michael on Delta
107 (7, 4, 'Sales Rep', 140), -- George on Delta
108 (10, 5, 'Lead Researcher', 190), -- Julia on Epsilon
109 (15, 5, 'Researcher', NULL); -- Oscar on Epsilon, NULL hours
110
111 -- Add an employee in a department that will be used for NOT IN examples
112 INSERT INTO departments (dept_name, location, monthly_budget) VALUES ('Intern Pool', '
    Remote', 10000.00);
113 INSERT INTO employees (emp_name, dept_id, salary, manager_id, performance_rating,
    last_bonus, hire_date) VALUES
114 ('Intern Zero', (SELECT dept_id FROM departments WHERE dept_name = 'Intern Pool'),
    20000.00, NULL, NULL, NULL, '2024-06-01');
115
116 -- Add a department with no employees for Exercise III.4
117 INSERT INTO departments (dept_name, location, monthly_budget) VALUES ('Empty Department'
    , 'Nowhere', 1000.00);

```

Listing 1: Dataset for Advanced WHERE Conditions Exercises

1 Practice Meanings, Values, Relations, Advantages, and Unique Uses

1.1 Subquery with IN

Problem: List the names and salaries of all employees who work in departments located in 'New York'.

1.2 Subquery with EXISTS

Problem: Find the names of departments that have at least one employee with a salary greater than \$85,000.

1.3 Subquery with ANY

Problem: List employees whose salary is greater than *any* salary in the 'Support' department. (This means their salary is greater than the minimum salary in the 'Support' department).

1.4 Subquery with ALL

Problem: Find employees in the 'Sales' department whose salary is less than *all* salaries in the 'Technology' department.

1.5 IS DISTINCT FROM

Problem: List employees whose `performance_rating` is different from 3. This list should include employees who have a NULL `performance_rating` (as NULL is distinct from 3).

1.6 IS NOT DISTINCT FROM

Problem: Find pairs of employees (display their names) who have the exact same `manager_id`, including cases where both employees have no manager (i.e., their `manager_id` is NULL). Avoid listing an employee paired with themselves.

2 Practice Disadvantages of All Its Technical Concepts

2.1 NOT IN with Subquery Returning NULL

Problem: Attempt to list employees who are NOT leads on any project using the condition `emp_id NOT IN (SELECT lead_emp_id FROM projects)`. Observe the (potentially unexpected) result given that `projects.lead_emp_id` can be NULL (e.g., 'NoLead Project'). Explain why this happens.

2.2 != ANY Misinterpretation

Problem: Find employees whose salary is not equal to *any* salary found in the 'Intern Pool' department. The 'Intern Pool' department currently has one employee ('Intern Zero') with a salary of \$20,000. Consider what happens if the 'Intern Pool' department had multiple distinct salaries (e.g., \$20,000, \$22,000). Explain the logical evaluation of `salary != ANY (subquery_salaries)` in such a scenario.

2.3 Performance of IS DISTINCT FROM vs. Standard Operators (Conceptual)

Problem: Consider finding employees where `performance_rating` is 3.

- Compare conceptually querying this using `performance_rating = 3` versus `performance_rating IS NOT DISTINCT FROM 3`.
- When might the `IS NOT DISTINCT FROM` approach be slightly less optimal if `performance_rating` is indexed and guaranteed NOT NULL? Discuss potential minor overheads or familiarity issues for developers.

2.4 Readability of EXISTS vs. IN for Simple Cases

Problem: Retrieve employees who are in departments listed in a small, explicit list of department IDs (e.g., department IDs 1 and 2).

- Write a query fragment using `IN` with a literal list.
- Write a query fragment using `EXISTS` with a subquery that provides these values.
- Compare the readability and conciseness of these two approaches for this specific simple case.

3 Practice Inefficient/Incorrect Alternatives vs. Advanced WHERE Conditions

3.1 Inefficient COUNT(*) vs. EXISTS for Existence Check

Problem: List department names that have at least one project associated with them (i.e., a project where `lead_emp_id` belongs to an employee in that department).

- **Task:** Write a query using an inefficient approach: a correlated subquery with `COUNT(*)` in the `WHERE` clause, checking if `COUNT(*) > 0`.
- **Consideration:** Why is using `EXISTS` generally more efficient for this type of "is there at least one?" check?

3.2 Verbose/Incorrect NULL Handling vs. IS DISTINCT FROM

Problem: Find employees whose `last_bonus` is not \$5000.00. This list should include employees whose `last_bonus` is `NULL`, as `NULL` is considered different from \$5000.00.

- **Task:** Write a query using a verbose approach: `(last_bonus <> 5000.00 OR last_bonus IS NULL)`.
- **Consideration:** How does `last_bonus IS DISTINCT FROM 5000.00` offer a more concise and less error-prone solution?

3.3 Complex NULL-aware Equality vs. IS NOT DISTINCT FROM

Problem: Find all employees whose `manager_id` is the same as Peter Pan's `manager_id`. (Peter Pan, `emp_id` 16, has a `NULL manager_id`). Do not include Peter Pan himself in the results.

- **Task:** Write a query using a complex approach: explicitly check for equality of `manager_id` with Peter Pan's `manager_id`, AND explicitly check if both the employee's `manager_id` and Peter Pan's `manager_id` are `NULL`.
- **Consideration:** How does using `IS NOT DISTINCT FROM` simplify this `NULL`-aware equality check and make the query more robust and readable?

3.4 Using LEFT JOIN and checking for NULL vs. NOT EXISTS

Problem: Find departments that have no employees. (Note: The dataset includes an 'Empty Department' specifically for this exercise).

- **Task:** Write a query using a common approach: `LEFT JOIN` the `departments` table with the `employees` table and then filter for departments where the employee's primary key (or any non-nullable employee column from the join) is `NULL`. Another variant could use `GROUP BY` and `HAVING COUNT(e.emp_id) = 0`.
- **Consideration:** For the specific task of checking non-existence, how does `NOT EXISTS` compare in terms of directness and potential efficiency?

4 Hardcore Problem Combining Previous Concepts

Problem Statement:

Identify "Key Departments" based on the following criteria:

1. The department name must contain either 'Tech' or 'HR' (case-sensitive as per standard SQL LIKE).
2. The department's `monthly_budget` IS NULL, OR its `monthly_budget` = 50000.00.
3. The department must have at least one "veteran" employee associated with it. This is determined by checking if there EXISTS such an employee in the department. A "veteran" employee is defined as someone:
 - Hired more than 8 years ago from `CURRENT_DATE`.
 - Whose `performance_rating` IS DISTINCT FROM 1 (i.e., their rating is not 1, or their rating is NULL).

For these "Key Departments", calculate the total `hours_assigned` to all their employees for projects that started on or after '2023-01-01'. If a department has no such employees or projects meeting this date criterion, their total hours should be displayed as 0.

Display the `dept_name` and the calculated `total_project_hours`. Order the results in descending order of `total_project_hours`, then by `dept_name` alphabetically. Limit the result to the top 3 departments.

Note for execution: To get more illustrative results for this specific hardcore problem, you might consider temporarily updating the Technology department's budget before running your solution:

```
1 -- UPDATE departments SET monthly_budget = 50000.00 WHERE dept_name = 'Technology';
```

Listing 2: Optional temporary data modification for Hardcore Problem

Remember to revert this change if you want to work with the original dataset for other exercises.

Previous Concepts to Integrate:

- **Advanced WHERE:** EXISTS, IS DISTINCT FROM.
- **Basic WHERE:** LIKE, OR, date comparisons (`>=`), arithmetic with dates (`CURRENT_DATE - INTERVAL`).
- **Intermediate SQL:**
 - Joins: `INNER JOIN`, `LEFT JOIN`.
 - Aggregators: `SUM()`, `GROUP BY`.
 - Null Space: `COALESCE`.
 - Date Functions: `CURRENT_DATE`, `INTERVAL`.
- **Ordering & Limiting:** `ORDER BY`, `LIMIT`.
- Consider using Common Table Expressions (CTEs) for clarity.