

Advanced SQL Query Techniques

Other Query Clauses & LATERAL Joins

Sequential SQL

May 16, 2025

Let's unlock some SQL magic for smarter data quests!

Contents

1	Other Query Clauses: OFFSET and FETCH	3
1.1	What Are They? (Meanings & Values)	3
1.2	Relations: How They Play with Others	3
1.3	How to Use Them: Structures & Syntax	4
1.4	Why Use Them? (Advantages)	5
1.5	Watch Out! (Disadvantages)	6
2	LATERAL Joins: The Row-by-Row Sidekick	7
2.1	What Are They? (Meanings & Values)	7
2.2	Relations: How They Play with Others	7
2.3	How to Use Them: Structures & Syntax	9
2.4	Why Use Them? (Advantages)	10
2.5	Watch Out! (Disadvantages)	11

1 Other Query Clauses: OFFSET and FETCH

PAGINATION and precise result slicing are common needs. Standard SQL offers OFFSET and FETCH for this, helping you grab just the data slice you need, no more, no less!

1.1 What Are They? (Meanings & Values)

💡 Core Concepts: OFFSET and FETCH

These clauses fine-tune which rows your query returns from an ordered set.

⦿ OFFSET *count* [ROW | ROWS]:

○ **Meaning:** This clause tells the database to skip the first *count* rows that would normally be returned by the query (after ORDER BY has done its job). It's like saying, "Start reading from page *count*+1."

○ **Value:** The *count* is a non-negative integer. If *count* is 0, no rows are skipped.

⦿ FETCH {FIRST | NEXT} [*count*] {ROW | ROWS} ONLY:

○ **Meaning:** This clause specifies that only the next *count* rows should be returned, after any OFFSET rows have been skipped. FIRST and NEXT are synonyms. If *count* is omitted, it defaults to 1. "Just give me this many records, please!"

○ **Value:** The *count* is a positive integer.

The **output** is a subset of rows from the potentially larger result set defined by the rest of the query.

1.2 Relations: How They Play with Others

OFFSET and FETCH don't live in a vacuum; they interact with clauses you already know.

💡 Synergies and Dependencies

Within "Other Query Clauses":

⦿ OFFSET and FETCH are like partners in pagination. You often use them together, but FETCH FIRST N ROWS ONLY can be used alone (like LIMIT N). Using OFFSET alone is less common but syntactically allowed; it implies you want all rows *after* the offset.

Relations with Previous SQL Concepts:

⦿ ORDER BY (**Basic SQL**): This is the most crucial relationship! OFFSET and FETCH operate on a result set whose order is **defined by** ORDER BY. Without ORDER BY, the notion of "first rows to skip" or "next rows to fetch" is ambiguous and can lead to unpredictable results. *Always sort your list before you slice it, or the pieces might surprise you!*

- Concepts like DESC, ordering by multiple columns, NULLS FIRST/LAST from Advanced ORDER BY are fully compatible and often necessary for stable pagination.
- ⦿ **LIMIT (Basic SQL):** In PostgreSQL and MySQL, LIMIT N OFFSET M is a common, non-standard shorthand. The SQL standard OFFSET M ROWS FETCH NEXT N ROWS ONLY is equivalent. FETCH FIRST N ROWS ONLY is the standard form of LIMIT N.
- ⦿ **WHERE Clause (Basic SQL):** Filtering with WHERE (including advanced conditions like subqueries in WHERE, BETWEEN, IN, LIKE) happens *before* ORDER BY, and thus before OFFSET and FETCH are applied. The pagination clauses work on the already filtered and sorted rows.
- ⦿ **Joins (INNER, LEFT, etc. from Intermediate & Complementary SQL):** Joins are resolved first to produce a combined result set. OFFSET and FETCH can then be applied to this joined, ordered result set.
- ⦿ **Subqueries in FROM clause (Complementary SQL):** You can apply OFFSET and FETCH to the result of a subquery. This is useful for paginating complex, pre-processed data.

```

1 SELECT *
2 FROM (
3     SELECT employeeId, salary, hireDate
4     FROM Employees
5     WHERE departmentId = 2 -- From Engineering
6     ORDER BY hireDate DESC -- Most recently hired first
7 ) AS RecentHires
8 OFFSET 5 ROWS
9 FETCH NEXT 10 ROWS ONLY; -- Get 6th to 15th recent hire

```

Listing 1: Paginating a subquery's result

- ⦿ **Set Operations (UNION, INTERSECT, EXCEPT from Complementary SQL):** OFFSET and FETCH can be applied to the final result of a set operation, provided an ORDER BY clause is also applied to the overall result of the set operation.

1.3 How to Use Them: Structures & Syntax

⌘ Syntax Blueprints: OFFSET and FETCH

These clauses are typically placed at the end of a SELECT statement, after the ORDER BY clause.

1. Using OFFSET only:

```

1 SELECT columnList FROM YourTable
2 ORDER BY sortColumn
3 OFFSET numberToSkip ROWS;

```

Note: Skips numberToSkip rows and returns all subsequent rows. Less common on its own for strict pagination.

2. Using FETCH only:

```

1 SELECT columnList FROM YourTable
2 ORDER BY sortColumn
3 FETCH FIRST numberOfRowsToGet ROWS ONLY;

```

Note: This is the SQL standard equivalent of `LIMIT numberOfRowsToGet`.

3. Using OFFSET and FETCH together (most common for pagination):

```

1 SELECT columnList FROM YourTable
2 ORDER BY sortColumn
3 OFFSET numberToSkip ROWS
4 FETCH NEXT numberOfRowsToGet ROWS ONLY;

```

Example: To get the 3rd page of 10 items per page, you'd use `OFFSET 20 ROWS FETCH NEXT 10 ROWS ONLY`.

Keywords:

- ⦿ ROW and ROWS are interchangeable.
- ⦿ FIRST and NEXT are interchangeable in the FETCH clause.

PostgreSQL Alternative (LIMIT / OFFSET): Many databases, including PostgreSQL, offer a more concise (but non-standard) syntax:

```

1 -- Equivalent to FETCH FIRST N ROWS ONLY:
2 SELECT columnList FROM YourTable ORDER BY sortColumn LIMIT N;
3
4 -- Equivalent to OFFSET M ROWS FETCH NEXT N ROWS ONLY:
5 SELECT columnList FROM YourTable ORDER BY sortColumn LIMIT N
   OFFSET M;

```

While convenient in PostgreSQL, pgAdmin4 also fully supports the standard OFFSET/FETCH syntax. For portability, the standard syntax is preferred.

🔗 Example: Finding the 2nd and 3rd most expensive products

Let's say we have a ProductSales table.

```

1 SELECT productName, unitPrice, saleDate
2 FROM ProductSales
3 ORDER BY unitPrice DESC, saleDate DESC -- Primary sort by price,
   secondary by date
4 OFFSET 1 ROW -- Skip the most expensive one
5 FETCH NEXT 2 ROWS ONLY; -- Get the next two (2nd and 3rd)

```

Listing 2: Fetching 2nd and 3rd priciest sales

This query skips the single highest-priced sale and then retrieves the next two.

1.4 Why Use Them? (Advantages)

👍 The Upsides of Slicing

- ⦿ **Efficient Pagination:** This is their superpower! They allow you to break down vast result sets into smaller, digestible "pages" for user interfaces or batch processing. *No need to fetch the whole phonebook just to find one number!*

- ⦿ **SQL Standard Compliance:** OFFSET and FETCH are part of the SQL standard (SQL:2008 and later), making your queries more portable across different database systems that adhere to the standard.
- ⦿ **Clarity of Intent:** The syntax clearly expresses the desire to skip a certain number of rows and then take a specific number.
- ⦿ **Top-N / Bottom-N Queries:** Easily retrieve the top or bottom N records according to some criteria when combined with ORDER BY ... ASC/DESC and FETCH FIRST N ROWS ONLY.
- ⦿ **Resource Management:** By fetching only necessary data, you reduce data transfer over the network and memory consumption on the client side.

1.5 Watch Out! (Disadvantages)

⚠ Potential Pitfalls

- ⦿ **Performance with Large OFFSET Values:** This is the big one. When you use a very large OFFSET (e.g., 'OFFSET 1000000 ROWS'), many database systems still need to internally generate, sort, and then count through all those million rows before they can fetch the few you asked for. This can be very slow. *Skipping a million pages in a book still means you touched them all!*
 - For "deep pagination," alternative techniques like keyset pagination (or "seek method") are often more performant. This involves using WHERE conditions on indexed columns from the last row of the previous page (e.g., 'WHERE (lastValue, id) > ('someVal', 123)'). (Keyset pagination uses concepts you already know but is a pattern, not a specific clause).
- ⦿ **Crucial Dependency on ORDER BY:** If you forget ORDER BY, or if your ORDER BY clause doesn't produce a unique, stable order, the rows skipped and fetched can be inconsistent across query executions. The database gives no guarantee on row order without ORDER BY. *A shifty list makes for shifty slices!*
- ⦿ **Non-Standard LIMIT Clause Prevalence:** Many developers are accustomed to the non-standard LIMIT offset, count or LIMIT count OFFSET offset syntax (e.g., MySQL, PostgreSQL, SQLite). While OFFSET/FETCH is standard, awareness of LIMIT variations is useful when working with different systems or older codebases.

2 LATERAL Joins: The Row-by-Row Sidekick

LATERAL joins are a powerful feature in SQL that allow a subquery in the FROM clause to reference columns from preceding tables in the same FROM clause. This enables complex, per-row calculations or lookups that are otherwise difficult or less efficient to express. Think of it as running a mini-query for each row of your main table!

2.1 What Are They? (Meanings & Values)

💡 Core Concept: LATERAL Joins

- 🕒 **Meaning:** A LATERAL join allows a derived table (subquery) or a table function on the right side of the join to be evaluated *for each row* from the table expression on its left. Critically, the right-hand side expression can **correlate** with columns from the left-hand side row currently being processed. It's like saying, "For this specific student, go find their top 3 recent test scores."
- 🕒 **Keywords:** The keyword LATERAL is used immediately before the subquery or table function in the FROM clause.
 - FROM LeftTable LT, LATERAL (subquery referencing LT.column) AS Sub
 - FROM LeftTable LT [INNER | LEFT | CROSS] JOIN LATERAL (subquery referencing LT.column) AS Sub ON ...
- 🕒 **Value/Output:** The result is a joined table.
 - If used with INNER JOIN LATERAL or CROSS JOIN LATERAL, rows from the left table are included only if the lateral subquery produces at least one row for them.
 - If used with LEFT JOIN LATERAL ... ON TRUE (or a suitable condition), all rows from the left table are included. If the lateral subquery produces no rows for a given left table row, columns from the lateral subquery will be NULL for that row.

The power comes from the subquery being re-evaluated dynamically for each outer row, using values from that outer row.

2.2 Relations: How They Play with Others

LATERAL joins unlock new ways to combine and process data by building upon familiar SQL concepts.

💡 Connections and Context

Within "LATERAL Joins" Concepts:

- 🕒 The LATERAL keyword itself is the enabler.
- 🕒 It's typically combined with standard join types like INNER JOIN LATERAL, LEFT

JOIN LATERAL, or simply a comma followed by LATERAL (which behaves like CROSS JOIN LATERAL or an implicit INNER JOIN LATERAL if the subquery returns no rows for an outer row).

Relations with Previous SQL Concepts:

- **Subqueries in FROM clause (Complementary SQL):** A LATERAL subquery *is* a subquery in the FROM clause. The key difference is its ability to reference preceding table columns. Regular subqueries in FROM cannot do this; they are evaluated independently. *LATERAL gives your subquery eyes to see its neighbors!*
- **Correlated Subqueries (Complementary SQL):** LATERAL brings the power of correlation, traditionally seen in scalar subqueries in the SELECT list or in WHERE clause predicates (like EXISTS), directly into the FROM clause. Unlike scalar correlated subqueries that must return a single value, a LATERAL subquery can return an entire set of rows (a table) for each outer row.
- **Joins (INNER, LEFT, CROSS - Intermediate & Complementary SQL):** LATERAL modifies the behavior of these joins. For example, LEFT JOIN LATERAL ensures all left-side rows are kept, even if the correlated subquery finds nothing for them.
- **WHERE Clause (Basic SQL):**
 - *Inside* the LATERAL subquery: This is where the correlation happens, filtering rows of the subquery based on values from the current outer row.
 - *Outside* the LATERAL join (in the main query's WHERE clause): Filters the combined result set produced by the LATERAL join.
- **ORDER BY / LIMIT (or OFFSET/FETCH) (Basic SQL & Other Query Clauses):** These are extremely powerful *inside* the LATERAL subquery. This is the standard way to solve "Top-N-per-group" problems (e.g., "for each department, get its top 3 highest-paid employees").
- **Aggregate Functions (SUM, AVG, etc. - Intermediate SQL):**
 - *Inside* the LATERAL subquery: Can compute aggregate values for each outer row (e.g., "for each product, calculate its total sales from another table").
 - *Outside*, in the main query: Can aggregate results from the LATERAL join.
- **Scalar Subqueries & Subqueries in SELECT (Complementary SQL):** LATERAL joins can often replace multiple, less efficient scalar correlated subqueries in the SELECT list, especially when you need several related pieces of information for each outer row.

2.3 How to Use Them: Structures & Syntax

Syntax Blueprints: LATERAL Joins

The LATERAL keyword is placed in the FROM clause.

1. Implicit Cross/Inner Join with LATERAL (Comma Syntax):

```
1 FROM LeftTableA A,  
2     LATERAL (  
3         SELECT SubColumn1, SubColumn2  
4         FROM SomeOtherTable SOT  
5         WHERE SOT.linkingColumn = A.columnFromA -- Correlation!  
6         -- ... other conditions, ORDER BY, LIMIT/FETCH ...  
7     ) AS SubqueryAlias
```

If the subquery returns no rows for a row in LeftTableA, that row from LeftTableA will NOT appear in the final result.

2. Explicit JOIN LATERAL (e.g., LEFT JOIN LATERAL):

```
1 FROM LeftTableA A  
2 LEFT JOIN LATERAL ( -- Could be INNER JOIN LATERAL, CROSS JOIN  
3     LATERAL  
4     SELECT SubColumn1, SubColumn2  
5     FROM SomeOtherTable SOT  
6     WHERE SOT.linkingColumn = A.columnFromA -- Correlation!  
7     -- ... other conditions, ORDER BY, LIMIT/FETCH ...  
8 ) AS SubqueryAlias ON TRUE -- ON TRUE is common for LEFT JOIN  
9                             LATERAL  
10                             -- when all logic is in the subquery.
```

With LEFT JOIN LATERAL, if the subquery returns no rows for a row in LeftTableA, that row from LeftTableA IS STILL included, with NULL values for columns from SubqueryAlias.

Key Use Case: Top-N per Group This is a classic LATERAL problem. Example: Get the latest 2 projects for each employee.

```
1 SELECT  
2     E.employeeId,  
3     E.firstName,  
4     E.lastName,  
5     P.projectName,  
6     P.assignmentDate  
7 FROM  
8     Employees E  
9 LEFT JOIN LATERAL (  
10     SELECT  
11         EP.projectName,  
12         EP.assignmentDate  
13     FROM  
14         EmployeeProjects EP  
15     WHERE  
16         EP.employeeId = E.employeeId -- Correlate to the current  
17         employee  
18     ORDER BY  
19         EP.assignmentDate DESC  
20     FETCH FIRST 2 ROWS ONLY -- Get the 2 most recent  
21 ) AS P ON TRUE -- ON TRUE as we want all employees  
22 ORDER BY
```

22

```
E.employeeId, P.assignmentDate DESC;
```

Listing 3: Top 2 recent projects per employee

This query shows each employee, and next to them, their two most recent projects. If an employee has fewer than two, it shows what they have. If none, project details are NULL.

Database Support:

- ⦿ **PostgreSQL:** Full support for LATERAL.
- ⦿ **MySQL:** Supported since version 8.0.14.
- ⦿ **Oracle:** Supports LATERAL.
- ⦿ **SQL Server:** Uses a similar concept called APPLY (CROSS APPLY and OUTER APPLY). CROSS APPLY is like an inner lateral join, and OUTER APPLY is like a left lateral join.

When working in pgAdmin4, you can use the standard LATERAL syntax.

2.4 Why Use Them? (Advantages)

The Power of Per-Row Processing

- ⦿ **Solves "Top-N-per-Group" Elegantly:** This is a common and often tricky problem (e.g., "find the 3 most recent orders for each customer"). LATERAL with ORDER BY and LIMIT/FETCH inside the subquery is the canonical SQL solution. *Like giving each group manager a tool to pick their own top performers!*
- ⦿ **Increased Readability for Complex Correlations:** For certain problems, LATERAL can make the query logic clearer than alternative solutions involving complex correlated scalar subqueries or convoluted self-joins. It explicitly states "for each row X, compute Y".
- ⦿ **Flexibility in Subqueries:** The lateral subquery can be arbitrarily complex, involving its own joins, aggregations, etc., all parameterized by the current outer row.
- ⦿ **Potential for Better Performance (sometimes):** Compared to multiple correlated scalar subqueries in the SELECT list (each hitting the database potentially), a single LATERAL join might be optimized better by the database engine, especially if the subquery can efficiently use indexes based on the correlation columns.
- ⦿ **Joining with Table-Valued Functions (TVFs):** LATERAL is essential for passing columns from an outer table as arguments to a TVF that returns a set of rows. (While specific Set Returning Functions like `generate_series` are later in the course, the general advantage of LATERAL with function-like constructs is significant).

2.5 Watch Out! (Disadvantages)

Points to Ponder

- ⦿ **Performance Considerations:** The lateral subquery is executed *for each row* of the outer table expression. If the outer table is large and the lateral subquery is complex or cannot use indexes efficiently, the query can become very slow. *Asking every person in a stadium for their life story takes time!*
 - Always ensure that the correlated conditions inside the LATERAL subquery can leverage indexes on the tables used within the subquery. Analyze query plans (EXPLAIN) for performance-critical LATERAL joins.
- ⦿ **Increased Complexity for Beginners:** The concept of a subquery being re-evaluated for each outer row can be less intuitive than simple joins for those new to the idea. *It's a step up from a simple handshake to a coordinated dance move.*
- ⦿ **Overkill for Simple Cases:** If a standard JOIN (e.g., INNER JOIN on a foreign key) can achieve the same result, using LATERAL would be unnecessarily complex and verbose. Stick to the simplest effective tool.
- ⦿ **Syntax Variations (e.g., APPLY in SQL Server):** While LATERAL is standardizing, if you work across different RDBMS, you might encounter vendor-specific alternatives like APPLY in SQL Server, which function similarly but have different syntax.

With OFFSET, FETCH, and LATERAL in your SQL toolkit, you're now equipped to craft even more precise and powerful data queries. Happy querying!