

Common Table Expressions - CTEs

Advanced Query Techniques: Exercises

May 18, 2025

Contents

1	Category (i): Practice Meanings, Values, Relations, and Advantages	6
2	Category (ii): Practice Disadvantages	7
3	Category (iii): Practice Cases Avoiding Inefficient Basic Solutions	8
4	Category (iv): Hardcore Combined Problem	9

Global Dataset for PostgreSQL

The following SQL code creates and populates the necessary tables for the exercises. Execute this script in your PostgreSQL environment before attempting the exercises.

```
1  -- Dataset for Category (i)
2  CREATE TABLE DepartmentsI (
3      departmentId INTEGER PRIMARY KEY,
4      departmentName VARCHAR(100) NOT NULL,
5      locationCity VARCHAR(50)
6  );
7
8  CREATE TABLE EmployeesI (
9      employeeId INTEGER PRIMARY KEY,
10     employeeName VARCHAR(100) NOT NULL,
11     departmentId INTEGER REFERENCES DepartmentsI(departmentId),
12     managerId INTEGER REFERENCES EmployeesI(employeeId), -- Self-reference for hierarchy
13     salary DECIMAL(10, 2) NOT NULL,
14     hireDate DATE NOT NULL
15 );
16
17 INSERT INTO DepartmentsI (departmentId, departmentName, locationCity) VALUES
18 (1, 'Technology', 'New York'),
19 (2, 'Human Resources', 'London'),
20 (3, 'Sales', 'Tokyo'),
21 (4, 'Marketing', 'Paris');
22
23 INSERT INTO EmployeesI (employeeId, employeeName, departmentId, managerId, salary,
24     hireDate) VALUES
25 (101, 'Alice Wonderland', 1, NULL, 120000.00, '2018-03-15'),
26 (102, 'Bob The Builder', 1, 101, 90000.00, '2019-07-01'),
27 (103, 'Charlie Brown', 1, 102, 80000.00, '2020-01-10'),
28 (104, 'Diana Prince', 2, NULL, 110000.00, '2017-05-20'),
29 (105, 'Eve Harrington', 2, 104, 75000.00, '2021-02-28'),
30 (106, 'Frankenstein Monster', 3, NULL, 130000.00, '2018-11-01'),
31 (107, 'Grace OMalley', 3, 106, 85000.00, '2019-05-15'),
32 (108, 'Henry Jekyll', 3, 106, 82000.00, '2022-08-20'),
33 (109, 'Ivy Pepper', 1, 101, 95000.00, '2020-06-01'),
34 (110, 'John Doe', NULL, 101, 60000.00, '2023-01-15');
35
36 -- Dataset for Category (ii)
37 CREATE TABLE ProductCategoriesII (
38     categoryId SERIAL PRIMARY KEY,
39     categoryName VARCHAR(50) UNIQUE NOT NULL
40 );
41
42 CREATE TABLE ProductsII (
43     productId SERIAL PRIMARY KEY,
44     productName VARCHAR(100) NOT NULL,
45     categoryId INTEGER REFERENCES ProductCategoriesII(categoryId),
46     basePrice DECIMAL(10,2)
47 );
48
49 CREATE TABLE SalesTransactionsII (
50     transactionId SERIAL PRIMARY KEY,
51     productId INTEGER REFERENCES ProductsII(productId),
52     saleDate TIMESTAMP NOT NULL,
53     quantitySold INTEGER NOT NULL,
54     discount DECIMAL(3,2) DEFAULT 0.00
55 );
56
57 INSERT INTO ProductCategoriesII (categoryName) VALUES ('Electronics'), ('Books'), ('Home
58     Goods');
59 INSERT INTO ProductsII (productName, categoryId, basePrice) VALUES
60 ('Laptop Pro', 1, 1200.00), ('Quantum Physics Primer', 2, 25.00), ('Smart LED Bulb', 3,
61     15.00),
62 ('Desktop Gamer', 1, 1800.00), ('History of Time', 2, 20.00), ('Robotic Vacuum', 3,
63     300.00);
64
65 DO $$
66 DECLARE
67     i INT;
```

```

64     pId INT;
65     sDate TIMESTAMP;
66     qty INT;
67 BEGIN
68     FOR i IN 1..10000 LOOP
69         pId := (MOD(i, 6)) + 1;
70         sDate := CURRENT_TIMESTAMP - (MOD(i,365) || ' days')::INTERVAL - (MOD(i,24) || '
            hours')::INTERVAL;
71         qty := (MOD(i, 5)) + 1;
72         INSERT INTO SalesTransactionsII (productId, saleDate, quantitySold, discount)
73         VALUES (pId, sDate, qty, CASE WHEN MOD(i,10) = 0 THEN 0.05 ELSE 0.00 END);
74     END LOOP;
75 END $$;
76
77 UPDATE SalesTransactionsII
78 SET saleDate = CURRENT_DATE - INTERVAL '1 month' + (MOD(transactionId, 30) || ' days')::
    INTERVAL
79 WHERE MOD(productId, 2) = 0; -- Update some products to have recent sales
80
81 -- Dataset for Category (iii)
82 CREATE TABLE CustomersIII (
83     customerId SERIAL PRIMARY KEY,
84     customerName VARCHAR(100) NOT NULL,
85     registrationDate DATE,
86     city VARCHAR(50)
87 );
88
89 CREATE TABLE ProductsMasterIII (
90     productId SERIAL PRIMARY KEY,
91     productName VARCHAR(100),
92     category VARCHAR(50)
93 );
94
95 CREATE TABLE OrdersIII (
96     orderId SERIAL PRIMARY KEY,
97     customerId INTEGER REFERENCES CustomersIII(customerId),
98     orderDate DATE,
99     shipmentRegion VARCHAR(50)
100 );
101
102 CREATE TABLE OrderItemsIII (
103     orderItemId SERIAL PRIMARY KEY,
104     orderId INTEGER REFERENCES OrdersIII(orderId),
105     productId INTEGER REFERENCES ProductsMasterIII(productId),
106     quantity INTEGER,
107     pricePerUnit DECIMAL(10,2)
108 );
109
110 INSERT INTO CustomersIII (customerName, registrationDate, city) VALUES
111 ('Global Corp', '2020-01-15', 'New York'), ('Local Biz', '2021-06-01', 'London'),
112 ('Alpha Inc', '2019-11-20', 'Tokyo'), ('Beta LLC', '2022-03-10', 'New York');
113
114 INSERT INTO ProductsMasterIII (productName, category) VALUES
115 ('Widget A', 'Gadgets'), ('Widget B', 'Gizmos'), ('Service C', 'Services'), ('Tool D', '
    Tools');
116
117 INSERT INTO OrdersIII (customerId, orderDate, shipmentRegion) VALUES
118 (1, '2022-02-10', 'North America'), (2, '2022-03-15', 'Europe'),
119 (1, '2023-04-20', 'North America'), (3, '2023-05-05', 'Asia'),
120 (2, '2023-06-10', 'Europe'), (4, '2022-07-01', 'North America');
121
122 INSERT INTO OrderItemsIII (orderId, productId, quantity, pricePerUnit) VALUES
123 (1, 1, 10, 50.00), (1, 2, 5, 100.00), (2, 3, 1, 200.00),
124 (3, 1, 20, 45.00), (3, 4, 2, 150.00), (4, 2, 8, 95.00),
125 (5, 3, 2, 190.00), (6, 4, 3, 140.00);
126
127 -- Dataset for Category (iv)
128 CREATE TABLE DepartmentsIV (
129     departmentId SERIAL PRIMARY KEY,
130     departmentName VARCHAR(100) NOT NULL,
131     headEmployeeId INTEGER -- Nullable, to be cross-referenced with EmployeesIV
132 );
133

```

```

134 CREATE TABLE EmployeesIV (
135     employeeId SERIAL PRIMARY KEY,
136     employeeName VARCHAR(100) NOT NULL,
137     departmentId INTEGER REFERENCES DepartmentsIV(departmentId),
138     managerId INTEGER REFERENCES EmployeesIV(employeeId), -- For hierarchy
139     salary DECIMAL(10, 2) NOT NULL,
140     hireDate DATE NOT NULL
141 );
142
143 ALTER TABLE DepartmentsIV ADD CONSTRAINT fkHeadEmployee FOREIGN KEY (headEmployeeId)
    REFERENCES EmployeesIV(employeeId) DEFERRABLE INITIALLY DEFERRED;
144
145 INSERT INTO DepartmentsIV (departmentId, departmentName) VALUES
146 (1, 'Engineering'), (2, 'Product Management'), (3, 'Research & Development'), (4, '
    Operations');
147
148 INSERT INTO EmployeesIV (employeeId, employeeName, departmentId, managerId, salary,
    hireDate) VALUES
149 (1, 'Ava CEO', 1, NULL, 250000, '2015-01-01'),
150 (2, 'Brian Lead', 1, 1, 150000, '2018-06-01'),
151 (3, 'Chloe SeniorDev', 1, 2, 110000, '2020-03-15'),
152 (4, 'David JuniorDev', 1, 3, 75000, '2022-07-01'),
153 (5, 'Eli PMHead', 2, 1, 160000, '2017-09-01'),
154 (6, 'Fiona SeniorPM', 2, 5, 120000, '2020-11-01'),
155 (7, 'George PM', 2, 6, 85000, '2021-05-10'),
156 (8, 'Hannah RDHead', 3, 1, 170000, '2016-04-12'),
157 (9, 'Ian SeniorScientist', 3, 8, 130000, '2021-01-20'),
158 (10, 'Julia Scientist', 3, 9, 90000, '2022-08-01'),
159 (11, 'Kevin OpsLead', 4, 1, 140000, '2019-02-10'),
160 (12, 'Liam OpsSpecialist', 4, 11, 95000, '2021-10-05'),
161 (13, 'Mike AnotherDev', 1, 2, 105000, '2021-02-01');
162
163 UPDATE DepartmentsIV SET headEmployeeId = 2 WHERE departmentName = 'Engineering';
164 UPDATE DepartmentsIV SET headEmployeeId = 5 WHERE departmentName = 'Product Management';
165 UPDATE DepartmentsIV SET headEmployeeId = 8 WHERE departmentName = 'Research &
    Development';
166 UPDATE DepartmentsIV SET headEmployeeId = 11 WHERE departmentName = 'Operations';
167
168 CREATE TABLE ProjectsIV (
169     projectId SERIAL PRIMARY KEY,
170     projectName VARCHAR(150) NOT NULL,
171     startDate DATE,
172     endDate DATE,
173     budget DECIMAL(12, 2)
174 );
175
176 CREATE TABLE TasksIV (
177     taskId SERIAL PRIMARY KEY,
178     projectId INTEGER REFERENCES ProjectsIV(projectId),
179     taskName VARCHAR(200),
180     assignedToEmployeeId INTEGER REFERENCES EmployeesIV(employeeId),
181     estimatedHours INTEGER,
182     actualHours INTEGER,
183     status VARCHAR(20)
184 );
185
186 CREATE TABLE TimeLogsIV (
187     logId SERIAL PRIMARY KEY,
188     taskId INTEGER REFERENCES TasksIV(taskId),
189     employeeId INTEGER REFERENCES EmployeesIV(employeeId),
190     logDate DATE NOT NULL,
191     hoursWorked DECIMAL(5,2) NOT NULL,
192     notes TEXT
193 );
194
195 INSERT INTO ProjectsIV (projectName, startDate, endDate, budget) VALUES
196 ('Alpha Core System', '2022-01-01', '2023-12-31', 200000.00),
197 ('Beta Mobile App', '2023-03-01', '2024-02-28', 80000.00),
198 ('Gamma Research Initiative', '2021-06-15', '2023-05-30', 160000.00),
199 ('Delta Operations Upgrade', '2023-07-01', NULL, 120000.00);
200
201 INSERT INTO TasksIV (projectId, taskName, assignedToEmployeeId, estimatedHours,
    actualHours, status) VALUES

```

```

202 (1, 'Design Alpha Architecture', 3, 100, 90, 'Completed'),
203 (1, 'Develop Alpha Module 1', 3, 150, 160, 'In Progress'),
204 (2, 'Beta UI/UX Design', 6, 80, 70, 'Completed'),
205 (2, 'Beta Backend Dev', 7, 120, 50, 'In Progress'),
206 (3, 'Gamma Initial Research', 9, 200, 180, 'Completed'),
207 (3, 'Gamma Experiment Setup', 9, 100, 110, 'Overdue'),
208 (4, 'Delta Process Analysis', 12, 60, 40, 'In Progress'),
209 (1, 'Alpha Documentation', 13, 80, 0, 'Pending');
210 -- The next task inserted will have taskId = 9 (due to SERIAL on previous 8 inserts)
211 INSERT INTO TasksIV (projectId, taskName, assignedToEmployeeId, estimatedHours,
    actualHours, status) VALUES
212 (2, 'Cross-project review for Alpha', 3, 20, 0, 'Pending');
213
214 INSERT INTO TimeLogsIV (logId, taskId, employeeId, logDate, hoursWorked, notes) VALUES
215 (DEFAULT, 1, 3, '2022-03-01', 8.0, 'Initial design'), (DEFAULT, 1, 3, '2022-03-02', 8.0,
    'Refinement'),
216 (DEFAULT, 2, 3, '2022-04-01', 8.0, 'Dev start'), (DEFAULT, 2, 3, '2022-04-02', 8.0, '
    Core logic'),
217 (DEFAULT, 3, 6, '2023-03-10', 7.0, 'UX flows'),
218 (DEFAULT, 5, 9, '2021-07-01', 6.0, 'Literature review'), (DEFAULT, 5, 9, '2021-07-02',
    8.0, 'Planning'),
219 (DEFAULT, 6, 9, '2021-09-01', 8.0, 'Setup phase 1'), (DEFAULT, 6, 9, '2021-09-02', 5.0,
    'Troubleshooting setup'),
220 (DEFAULT, 7, 12, '2023-07-15', 8.0, 'Mapping current state'),
221 (DEFAULT, 8, 13, '2022-05-01', 4.0, 'Doc outline');
222 -- logId values will be 1 to 11 after these inserts
223 INSERT INTO TimeLogsIV (logId, taskId, employeeId, logDate, hoursWorked, notes) VALUES
224 (DEFAULT, 2, 3, '2022-04-03', 8.0, 'Task 2 for emp 3'), -- emp 3 (Chloe) on task 2 (
    project 1), logId 12
225 (DEFAULT, 4, 7, '2023-08-01', 5.0, 'Task 4 for emp 7'), -- emp 7 (George) on task 4 (
    project 2), logId 13
226 (DEFAULT, 9, 3, '2023-09-01', 3.0, 'Time for task 9, project 2'); -- emp 3 works on task
    9 (project 2), logId 14

```

Listing 1: Global Dataset for Exercises

1 Category (i): Practice Meanings, Values, Relations, and Advantages

These exercises focus on understanding the fundamental meanings, values, and relational aspects of Common Table Expressions (CTEs). They demonstrate unique uses and advantages, building upon concepts from Basic and Intermediate SQL.

1. Exercise 1: Basic CTE for Readability

Problem: List all employees in the 'Technology' department who earn more than \$90,000. Show how a CTE can simplify selecting the department first.

2. Exercise 2: CTE Referenced Multiple Times

Problem: Find all employees whose salary is above the average salary of their respective department. Also, show the department's average salary. This requires calculating departmental average salary and then using it for comparison.

3. Exercise 3: Nested CTEs

Problem: List employees from 'New York' or 'London' who were hired after 2019. First, create a CTE for relevant departments. Then, a CTE for employees in those departments hired after 2019.

4. Exercise 4: Recursive CTE for Hierarchical Data

Problem: Display the organizational hierarchy for 'Charlie Brown' (employeeId 103), showing his reporting line up to the top manager. List employee ID, name, manager ID, and level in hierarchy.

2 Category (ii): Practice Disadvantages

These exercises explore potential disadvantages or limitations associated with CTEs, such as performance considerations and scope.

1. Exercise 1: Potential Performance Issue (Optimization Fence / Materialization)

Problem: Calculate the total revenue (price * quantity * (1-discount)) for each product using tables from dataset part II ('ProductsII', 'SalesTransactionsII', 'ProductCategoriesII'). Then, retrieve this information ONLY for products in the 'Electronics' category. A CTE might calculate revenue for ALL products first, then filter. (This exercise highlights a *potential* disadvantage; actual performance depends on the DBMS optimizer).

2. Exercise 2: No Indexing on CTE Results

Problem: Using tables from dataset part II, identify products that had sales in the month immediately preceding the current month (e.g., if today is Feb 15, 2024, identify sales in Jan 2024). Simulate multiple conceptual uses of this intermediate result: first list the product names, then provide a count of these distinct products. The disadvantage illustrated is that if the CTE result was large and queried multiple times, it's re-evaluated or its unindexed materialized result is scanned.

3. Exercise 3: CTE Scope Limitation

Problem: You need to calculate total sales revenue for the 'Books' category (using dataset part II tables) and use this total in two *separate subsequent independent queries* (e.g., one to show the total, another to show 10% of this total). Show conceptually or by attempting that a CTE defined in one query is not available in the next, illustrating its scope. Then, demonstrate how you would achieve this by re-declaring the CTE if needed.

3 Category (iii): Practice Cases Avoiding Inefficient Basic Solutions

These exercises demonstrate scenarios where CTEs offer significant advantages in readability, maintainability, and sometimes performance over more basic or convoluted approaches that don't leverage CTEs effectively. Use tables from dataset part III ('CustomersIII', 'ProductsMasterIII', 'OrdersIII', 'OrderItemsIII').

1. **Exercise 1: Replacing Repeated Subqueries**

Problem: Find customers who placed orders in both 2022 and 2023. List their names and city. Illustrate how CTEs can avoid repeating subquery logic that might scan 'OrdersIII' multiple times.

2. **Exercise 2: Simplifying Complex Joins and Filters**

Problem: List products (name and category) that were part of orders shipped to 'North America' and had a total order value (sum of quantity * pricePerUnit for all items in that order) greater than \$600. Show how CTEs can break down this logic compared to a single, very long query.

3. **Exercise 3: Avoiding Temporary Tables for Single-Query Scope**

Problem: Calculate the average total order value for each 'shipmentRegion'. Then, list regions whose average order value is greater than the overall average order value across all regions. Demonstrate how CTEs provide a cleaner, single-statement solution compared to potentially using temporary tables.

4. **Exercise 4: Step-by-Step Multi-Level Aggregations (Revised)**

Problem: For each product category, find the month (e.g., '2023-04') with the highest total sales quantity for that category. Display category, the best month, and total quantity for that month. Solve this using CTEs for structured aggregation, without using window functions.

4 Category (iv): Hardcore Combined Problem

This problem requires combining various SQL concepts learned prior to and including Common Table Expressions, focusing on their application in a complex scenario. Use tables from dataset part IV ('DepartmentsIV', 'EmployeesIV', 'ProjectsIV', 'TasksIV', 'TimeLogsIV'). Window functions (like 'RANK()', 'ROW_NUMBER() OVER()') should NOT be used as they are covered later in the course sequence.

1. Hardcore Problem

Problem: Identify the top 2 departments by the total salary of their 'Senior' employees. A 'Senior' employee is defined as someone with a salary > \$70,000 AND hired on or after '2020-01-01' AND has logged time (in 'TimeLogsIV') on at least one task belonging to a 'Critical' project. A 'Critical' project is any project from 'ProjectsIV' with a budget > \$150,000.

For these top 2 departments, display:

- a. Department Name.
- b. Total salary of these qualified 'Senior' employees in that department.
- c. The count of such 'Senior' employees in the department.
- d. Using a LATERAL join, for each of these top 2 departments, find the employee (can be any employee in that department, not necessarily senior) who has logged time against the highest number of distinct projects (based on 'TimeLogsIV' and 'TasksIV'). If there's a tie in distinct project count, pick the one with the lower 'employeeId'. Show this employee's name and their distinct project count.

Additionally, for the single department (from the top 2 identified above) that has the absolute highest total senior salary:

- e. Display the organizational hierarchy for its department head (employee specified in 'DepartmentsIV.headEmployeeId'), showing the reporting line upwards to the CEO (employee with 'managerId IS NULL'). List employee ID, name, manager ID, and level in hierarchy (0 for the department head, increasing for their managers).

Constraints:

- Departments must have at least 2 qualified 'Senior' employees to be considered for the top 2.
- Use 'FETCH FIRST ... ROWS ONLY' to get the top 2 departments based on total senior salary (descending).
- The final list of top 2 departments should be ordered by their total senior salary in descending order.
- The hierarchy should be for the head of the #1 department from this list.
- All parts of the problem should be solved within a single SQL statement where possible (the hierarchy query might be separate if needed for clarity, but aim to use CTEs effectively if combining). Ideally, two main 'SELECT' statements: one for the top 2 departments' details, and one for the hierarchy of the #1 department's head.